

Predicting Review Ratings from Yelp Dataset

Christian Becerra
Nicholas Hansen
Kheva Mann



Problem Statement

- Want to predict the rating of reviews
- Predict ratings given by users to businesses based on reviews most textually similar to it
- Ratings are given on a 1 to 5 scale
- Reviews come from a very large (6 GB) dataset from the company Yelp
- Publicly available and intended for academic use





Solution

- Use Natural Language Processing with the Spark framework on the Hadoop FS
- Reviews are converted into term frequency-inverse document frequency vectors or TF-IDF
- Cosine similarity is then detected
- K-nearest neighbor algorithms allow us to determine our final prediction
- This technique is fairly simple but sufficient for our needs
- Definitely has room for improvement




Too much data

- The initial data set from Yelp is approximately 6 GB
 - Much too large, especially for creating TF-IDF vectors
- Data must be reduced
 - Only a subset of data was taken (less than 10 MB)
 - Remaining data was filtered to remove unnecessary information
 - Calculating similarity amongst all documents explodes data size to n^2

Cleaning and Vectorizing the Text

- Review text must be vectorized (made numerical) in order to calculate similarity between documents
- Don't want noise in our data
 - Remove common stop words like “a” and “the” that provide little information (modified to keep common sentiment words)
 - Remove punctuation and extra spacing
 - Did not stem or expand contractions (possible improvement)



https://www.researchgate.net/figure/Dimensional-reduction-of-term-frequency-inverse-document-frequency-TF-IDF_fig2_330121599



TF-IDF Implementation

- TF-IDF determines how significant a word presence in a document is based upon the frequency of that word in the corpus.
 - Good for determining document uniqueness and clustering
 - $\text{TF-IDF}(t, d, D) = \text{TF}(t, d) * \text{IDF}(t, D)$
 - t = term
 - d = document
 - D = corpus of documents
- Considered top 1000 most prevalent words (better for sentiment)
- Used Spark MLlib sparse vectors to represent (TF-IDF vectors are cluttered with 0s)

$$w_{i,j} = tf_{i,j} \times \log \left(\frac{N}{df_i} \right)$$

tf_{ij} = number of occurrences of i in j

df_i = number of documents containing i

N = total number of documents

<https://medium.com/@imamun/creating-a-tf-idf-in-python-e43f05e4d424>



TF-IDF Implementation (Continued)

- TF-IDF normalized based on magnitudes allowing for fair comparisons
- Dot product is computed between all pairs of normalized document TF-IDF vectors (except self comparisons)
 - Equivalent to cosine similarity
 - Used to find most similar documents



KNN Implementation

- Regressive vs Classifier
 - Regressive
 - Average rating of the k nearest neighbors to make prediction
 - Useful for continuous variables (value predictions \$)
 - Classifier
 - Plurality rating of the k nearest neighbors classifies variable into predefined class
 - Useful for categorical variables (interest predictions)
- Did both!



Results (Methods)

- Tested with 500, 1000, and 1200 documents per rating respectively
- 500 reviews per rating worst performer
- Constant 1000 top words TF-IDF vector size
- Tested both regressive and classifier KNN
 - $K = \sqrt{\text{total_documents}}$
 - Standard K value



Results (1000 vs 1200) - Classifier

- Performed relatively similar
 - 1000 better at predicting 2.0 rating
 - 1200 better at predicting 4.0 rating
 - Likely transient rather than improved performance

1000 Reviews Per Rating

Overall Accuracy: 44.52%

F1-Score(1.0) = 0.5946428571428573

F1-Score(2.0) = 0.36412459720730395

F1-Score(3.0) = 0.330672748004561

F1-Score(4.0) = 0.331145885916204

F1-Score(5.0) = 0.5575589459084604

1200 Reviews Per Rating

Overall Accuracy: 44.55%

F1-Score(1.0) = 0.5947521865889213

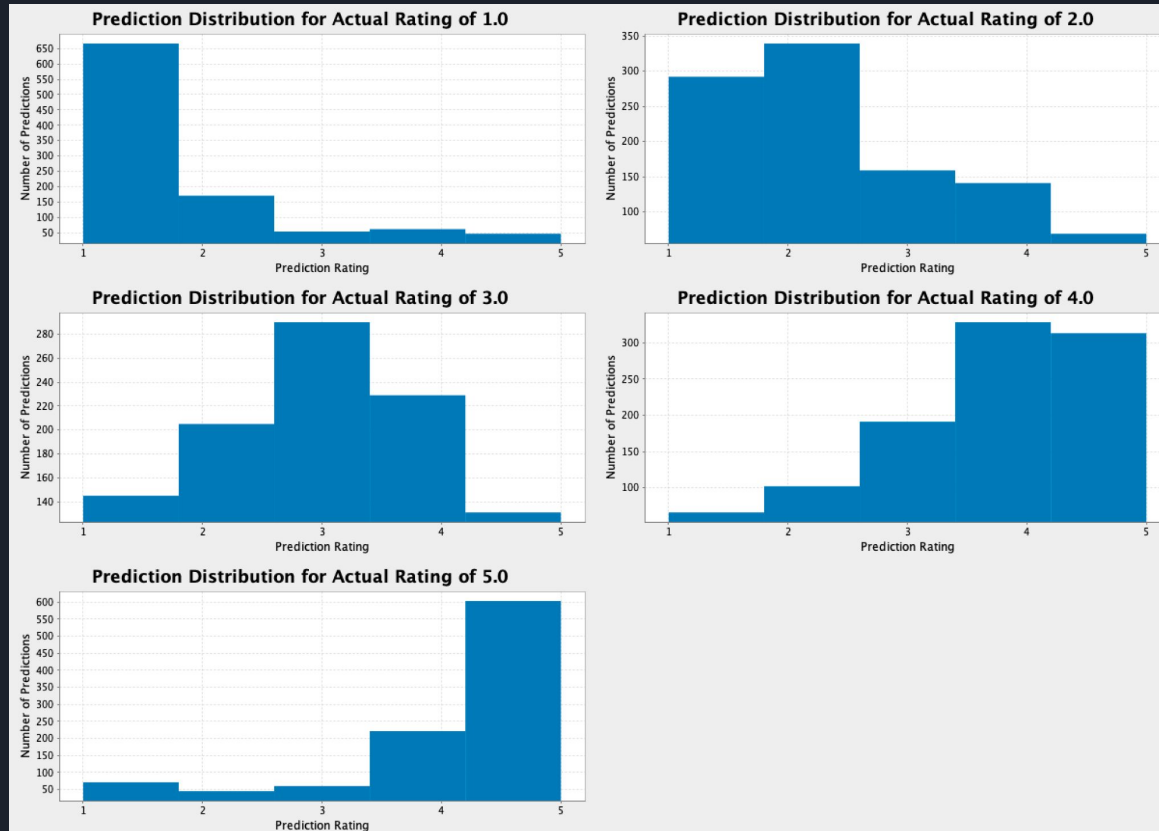
F1-Score(2.0) = 0.34399629114510893

F1-Score(3.0) = 0.33113828786453436


F1-Score(4.0) = 0.35070306038047977

F1-Score(5.0) = 0.5557729941291585

Results (1000 per Rating) - Classifier



Results (1200 per Rating) - Classifier





Results (1000 vs 1200) - Regressor

- Relatively the same

1000 Reviews Per Rating

RMSE = 1.137 stars

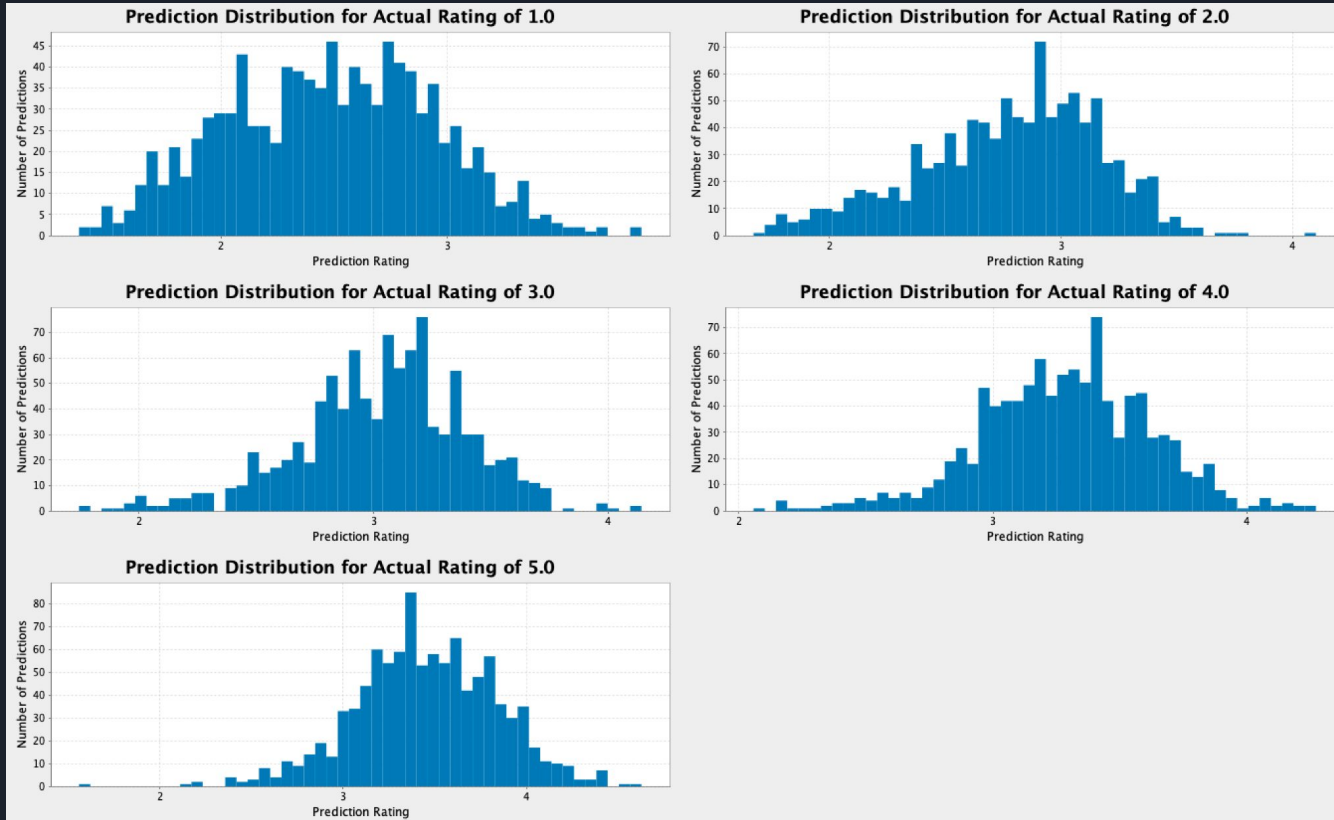
MAE = 0.964 stars

1200 Reviews Per Rating






RMSE = 1.132 stars

MAE = 0.96 stars

Results (1000 per Rating) - Regressor


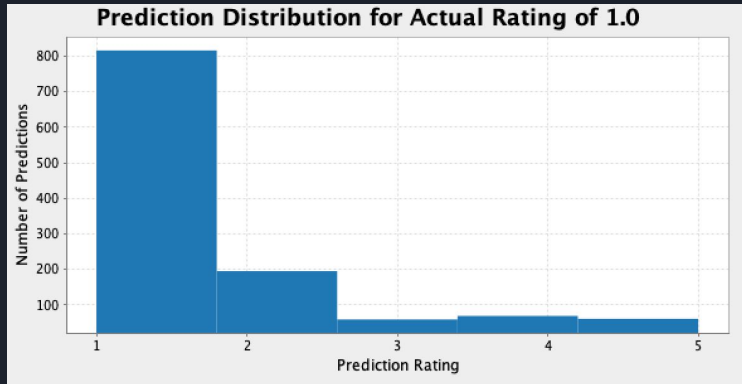


Results (1200 per Rating) - Regressor



Results: Classifier vs Regressor KNN

- Classifier good at predicting values at the bounds (1 star and 5 stars)
 - Decent at predicting values in the center (2-4 stars)
- Regressor decent at predicting values in the center
 - Bad at predicting values at the bounds
 - Taking the average biases values towards the center





What We Learned

- Vectorizing large amounts of data taking too long
 - Had to use Spark Mllib library for Sparse Vectors
- Learned a lot about text vectorization
- Learned about natural language processing and sentiment analysis
- Learned about KNN and two different ways of applying it
- Data visualization using MLlibrary
- Rudimentary due to methodology used compared to other strategies

Thank You!

